

# Linear Regression with multiple variables

## Multiple features

### Multiple features (variables).

Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
$x_1$	$x_2$	$x_3$	$x_4$	$y$
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...	...	...	...	...

Notation:

- $\rightarrow n$  = number of features  $n=4$
- $\rightarrow x^{(i)}$  = input (features) of  $i^{th}$  training example.
- $\rightarrow x_j^{(i)}$  = value of feature  $j$  in  $i^{th}$  training example.

$m=47$

$x^{(2)} = \begin{bmatrix} 1416 \\ 3 \\ 2 \\ 40 \end{bmatrix}$

$x_3^{(2)} = 2$

Andrew Ng

Multivariate linear regression:

$$\rightarrow h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

For convenience of notation, define  $x_0 = 1$ . ( $x_0^{(i)} = 1$ )

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n$$

$$= \theta^T x$$

$\theta^T$  is a  $(n+1) \times 1$  matrix.

## Gradient descent for multiple variables

Hypothesis:  $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$

Parameters:  $\theta_0, \theta_1, \dots, \theta_n$   $\theta$   $n+1$ -dimensional vector

Cost function:

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$J(\theta)$

Gradient descent:

Repeat {

$\rightarrow \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$

} (simultaneously update for every  $j = 0, \dots, n$ )

**Gradient Descent**

Previously ( $n=1$ ):

Repeat {

→  $\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$

$\frac{\partial}{\partial \theta_0} J(\theta)$

→  $\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$

(simultaneously update  $\theta_0, \theta_1$ )

}

**New algorithm ( $n \geq 1$ ):**

Repeat {

→  $\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$

(simultaneously update  $\theta_j$  for  $j = 0, \dots, n$ )

}

→  $\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$

→  $\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$

→  $\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)}$

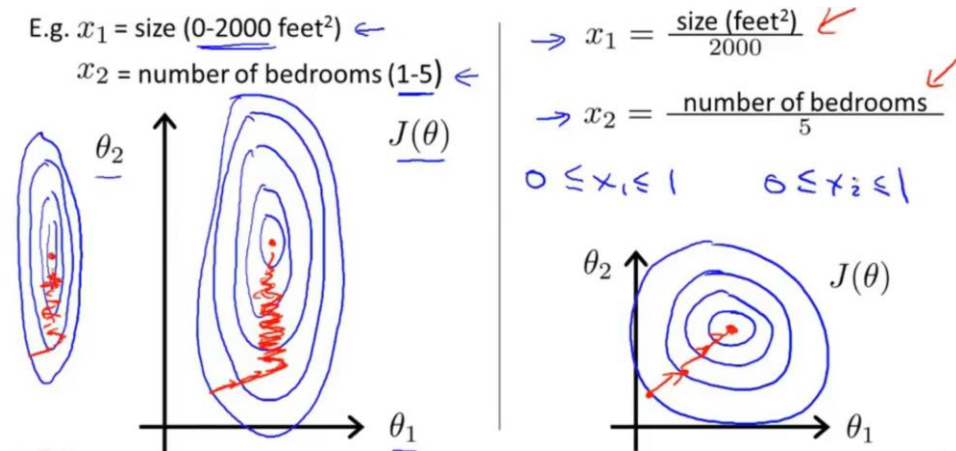
...

Andrew Ng

## Gradient descent in practice I: Feature Scaling

Feature Scaling

Idea: Make sure features are on a similar scale. (Then gradient descents can converge more quickly)



Get every feature into approximately a  $-1 \leq x_i \leq 1$  range.

$x_0 = 1$

$0 \leq x_1 \leq 3$  ✓

$-2 \leq x_2 \leq 0.5$  ✓

$-100 \leq x_3$  100 ✗

$-0.0001 \leq x_4$  0.0001 ✗

↑

$-3$  to  $3$  ✓

$-\frac{1}{3}$  to  $\frac{1}{3}$  ✓

## Mean normalization

Replace  $x_i$  with  $x_i - \mu_i$  to make features have approximately zero mean (Do not apply to  $x_0 = 1$ ).

E.g.  $x_1 = \frac{\text{size} - 1000}{2000}$       Average size = 1000  
 $x_2 = \frac{\# \text{bedrooms} - 2}{5}$       1-5 bedrooms

$-0.5 \leq x_1 \leq 0.5$        $-0.5 \leq x_2 \leq 0.5$

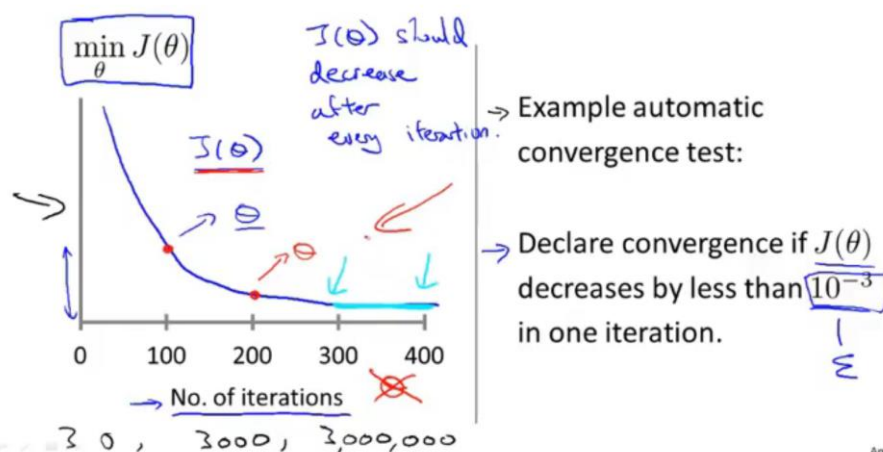
$x_1 \leftarrow \frac{x_1 - \mu_1}{s_1}$        $x_2 \leftarrow \frac{x_2 - \mu_2}{s_2}$

← avg. value of  $x_1$  in training set  
 ← range (max - min) (or standard deviation)

Andrew Ng

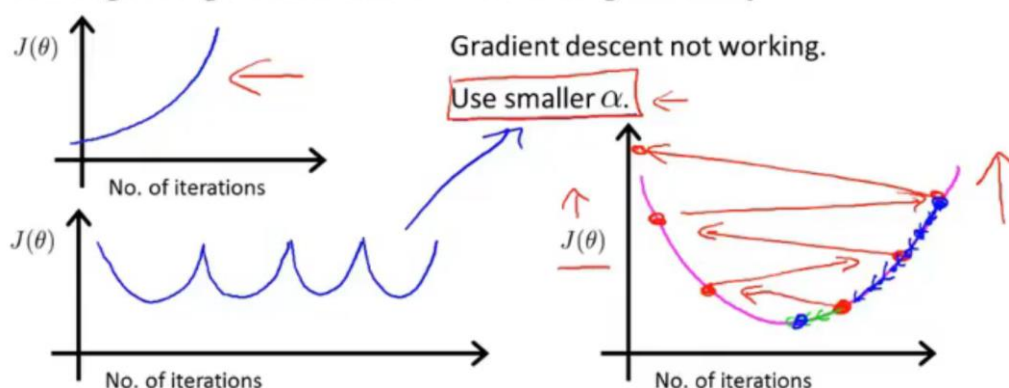
## Gradient descent in practice II: Learning rate

Making sure gradient descent is working correctly.



Andrew Ng

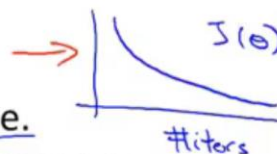
Making sure gradient descent is working correctly.



- For sufficiently small  $\alpha$ ,  $J(\theta)$  should decrease on every iteration.
- But if  $\alpha$  is too small, gradient descent can be slow to converge.

## Summary:

- If  $\alpha$  is too small: slow convergence.
- If  $\alpha$  is too large:  $J(\theta)$  may not decrease on every iteration; may not converge. (Slow converge also possible)



To choose  $\alpha$ , try

..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, ...

Arrows indicate a sequence of values, with some values being approximately 3 times the previous one (e.g., 0.003 is ~3x 0.001, 0.01 is ~3x 0.003, etc.).

Andrew Ng

## Feature and polynomial regression

define new features:  $x = \text{frontage} \times \text{depth}$ .

### Housing prices prediction

$$h_{\theta}(x) = \theta_0 + \theta_1 \times \underbrace{\text{frontage}}_{x_1} + \theta_2 \times \underbrace{\text{depth}}_{x_2}$$



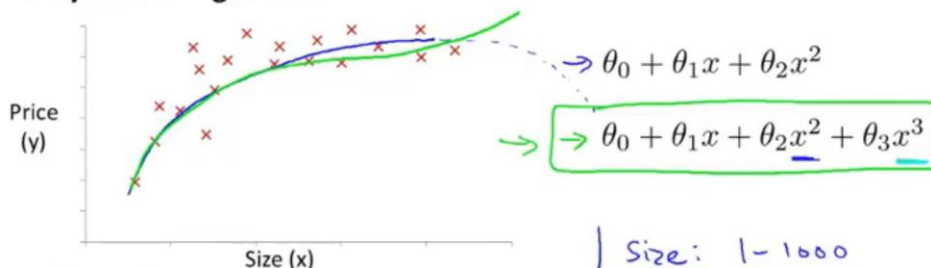
Area

$$x = \text{frontage} \times \text{depth}$$

$$h_{\theta}(x) \approx \theta_0 + \theta_1 x$$

land area

### Polynomial regression



$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

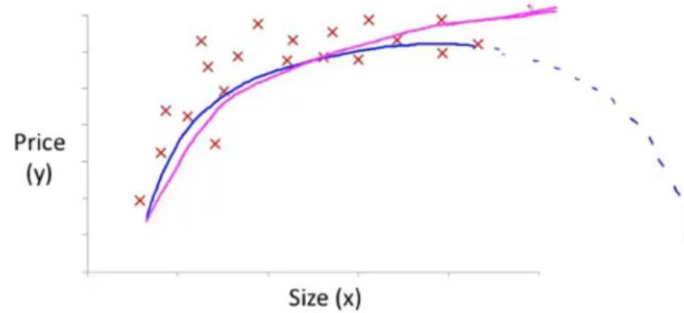
$$= \theta_0 + \theta_1 (\text{size}) + \theta_2 (\text{size})^2 + \theta_3 (\text{size})^3$$

$$\begin{aligned} \rightarrow x_1 &= (\text{size}) \\ \rightarrow x_2 &= (\text{size})^2 \\ \rightarrow x_3 &= (\text{size})^3 \end{aligned}$$

Size: 1 - 1000  
Size<sup>2</sup>: 1 - 1,000,000  
Size<sup>3</sup>: 1 - 10<sup>9</sup>

Andrew Ng

## Choice of features



$$\rightarrow h_{\theta}(x) = \theta_0 + \theta_1(\text{size}) + \theta_2(\text{size})^2$$

$$\rightarrow h_{\theta}(x) = \theta_0 + \theta_1(\text{size}) + \theta_2\sqrt{(\text{size})}$$

Andrew Ng

## Normal equation

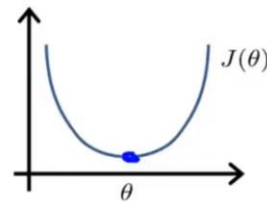
Normal equation: Method to solve for theta analytically.

Intuition: If 1D ( $\theta \in \mathbb{R}$ )

$$\rightarrow J(\theta) = a\theta^2 + b\theta + c$$

$$\frac{\partial}{\partial \theta} J(\theta) = \dots \stackrel{\text{set}}{=} 0$$

Solve for  $\theta$



$$\theta \in \mathbb{R}^{n+1} \quad J(\theta_0, \theta_1, \dots, \theta_m) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \dots \stackrel{\text{set}}{=} 0 \quad (\text{for every } j)$$

Solve for  $\theta_0, \theta_1, \dots, \theta_n$

Examples:  $m = 4$ .

	Size (feet <sup>2</sup> )	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
$x_0$	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}$$

$m \times (n+1)$

$$y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

$m$ -dimensional vector

$$\theta = (X^T X)^{-1} X^T y$$



$m$  examples  $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$ ;  $n$  features.

$$\underline{x^{(i)}} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ x_2^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \in \mathbb{R}^{n+1} \quad \left| \quad \begin{matrix} \text{X} \\ \text{(design matrix)} \end{matrix} \right. = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix}$$

E.g. If  $\underline{x^{(i)}} = \begin{bmatrix} 1 \\ x_1^{(i)} \end{bmatrix}$   $\rightarrow$   $\underline{X} = \begin{bmatrix} 1 & x_1^{(1)} \\ 1 & x_1^{(2)} \\ \vdots & \vdots \\ 1 & x_1^{(m)} \end{bmatrix}$   $\left| \quad \underline{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$

$\Theta = (X^T X)^{-1} X^T y$  (m x 2) n x (n+1)

Andrew Ng

Feature scaling isn't actually necessary in normal equation.

$$\underline{\theta} = \underline{(X^T X)^{-1} X^T y}$$

$(X^T X)^{-1}$  is inverse of matrix  $\underline{X^T X}$ .

Set  $\underline{A} = \underline{X^T X}$

$$\underline{(X^T X)^{-1}} = \underline{A^{-1}}$$

Octave: `pinv(X' * X) * X' * y`

$$\underline{\theta} = \underline{(X^T X)^{-1} X^T y}$$

$\min J(\theta)$

$\underline{X'}$   $\underline{X^T}$

~~Feature Scaling~~

$0 \leq x_1 \leq 1$

$0 \leq x_2 \leq 1000$

$0 \leq x_3 \leq 10^{-5}$  ✓

$m$  training examples,  $n$  features.

Gradient Descent

- • Need to choose  $\alpha$ .
- • Needs many iterations.
- Works well even when  $n$  is large.

$\rightarrow$

$\underline{n = 10^6}$

Normal Equation

- • No need to choose  $\alpha$ .
- • Don't need to iterate.
- Need to compute
- $\underline{(X^T X)^{-1}}$  n x n  $O(n^3)$
- Slow if  $n$  is very large.

$n = 100$

$n = 1000$

$n = 10000$

← -

## Normal equation and non-invertibility

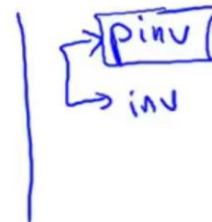
### Normal equation

$$\theta = (X^T X)^{-1} X^T y$$

$$\underline{X^T X}$$

- What if  $X^T X$  is non-invertible? (singular/degenerate)

- Octave:  $\text{pinv}(X' * X) * X' * y$



What if  $X^T X$  is non-invertible?

- Redundant features (linearly dependent).

E.g.  $x_1 = \text{size in feet}^2$

$$1 \text{ m} = 3.28 \text{ feet}$$

~~$x_2 = \text{size in m}^2$~~

$$x_1 = (3.28)^2 x_2$$

$$\rightarrow m = 10 \leftarrow$$

$$\rightarrow n = 100 \leftarrow$$

$$\theta \in \mathbb{R}^{101}$$

- Too many features (e.g.  $m \leq n$ ).

- Delete some features, or use regularization.

↓ later